

CONTEXT SHIFTING ARGUMENTS¹

Herman Cappelen
Vassar College

Ernie Lepore
Rutgers University

Context Shifting Arguments (CSA) ask us to consider two utterances of an unambiguous, non-vague, non-elliptic sentence S. If the consensus intuition is that what's said or expressed, or the truth-conditions and so possibly the truth-values, of these utterances differ, then CSA concludes S is context sensitive. Consider, for example, simultaneous utterances of 'I am wearing a hat', one by Stephen, one by Jason. Intuitively, these utterances might disagree in truth-value, contingent upon who is or isn't wearing a hat, and so they express distinct propositions and differ in their truth conditions. Since these differences are not the result of ambiguity (lexical or structural), vagueness, conversational implicature, or syntactic ellipsis, we have pretty strong evidence that 'I am wearing a hat' is context sensitive.

A central premise in all CSA is the assumption that the intuitions appealed to (about various utterances of S) are semantic; *viz.*, they are the sorts of intuitions a semantic theory must accommodate. One goal of this paper is to present tests for the semantic significance of such intuitions. If we are right, unsound CSA are ubiquitous. As a consequence a wide range of philosophers have mistaken views about context sensitivity.

Our attack on these positions centers around three closely related tests for context sensitivity.

- 1) We show that an expression e is context sensitive just in case it passes what we call the *Inter-Contextual Disquotation* test and an expression passes this test just in case there can be what we call a *Real CSA* for it. Ordinary context sensitive expressions, e.g., 'I', 'she', and 'now', pass this test (and there can be Real CSA for them), whereas, we will argue, controversial cases, such as 'know', 'good', and 'red' don't (and there are no Real CSA for them).

- 2) An expression *e* is context sensitive just in case it, in general, blocks what we call *inter-contextual disquotational indirect reports* (this is the IDI-Test). The controversial cases don't pass this test.
- 3) An expression *e* is context sensitive just in case it blocks what we call *Collective Descriptions*. Again, we argue that genuine context sensitive expressions block collective descriptions, whereas the controversial cases do not.

We consider a possible reply according to which the controversial cases fail to pass various tests for context sensitivity because of what can be called *contextual salience absorption*. This response, we show, fails.

At the end we present two possible diagnoses of how one can get fooled into thinking there is semantic context sensitivity when there is none.

At the very end of the paper there's an appendix about comparative adjectives. They don't seem to pass our tests. We outline some strategies for dealing with their apparent failure.

Contextualism and Context Shifting Arguments

Contextualism (as we use this expression) is the view that our language contains non-obvious contextually sensitive expressions. David Lewis (1979; 1996), among others,² defends an *epistemic* contextualism according to which ordinary knowledge-attributions, for example (1), are context sensitive.

- (1) Lewis knows that penguins eat fish.

Lewis invokes CSA in arguing that distinct utterances of sentences like (1) can differ in truth-value contingent upon which epistemic standards are summoned. In a philosophical context an utterance of (1) might be (deemed) false; but in ordinary non-philosophical ones it might not be.

Unger defends an *ethical* contextualism arguing that 'in many cases, the truth value...of a judgment about whether a person's behavior is morally permissible depends on the context in which that judgment is made or is grasped' (Unger 1995, p. 2).³ In defense of ethical contextualism, CSA is again invoked: 'For one thing, the content of a moral claim or belief is, on my view, relative to a context. For another (and this is really just a consequence of the first) two people in different contexts may utter 'x is good' and 'x is not good' and both speak truly' (Dreier 1990, p. 7).

Contextualism about quantifier expressions has also become rather widespread (e.g., see, Westerståhl 1989, Stanley/Szabo 2000, Stanley 2002; 2003). Intuition has it that simultaneous utterances of (2) might also disagree in truth-value, contingent on which domain of quantification is contextually salient.

- (2) Everyone took the exam.

If the domain includes everyone registered for a certain class, an utterance of (2) might be (reckoned as) true; but in another context where it might include the entire student body an utterance of (2) might be (considered) false. (2) does not change meaning across contexts of use, but the contextually determined domain does.⁴

Radical contextualism is the view that *every single linguistic expression* is subject to context shifting. So consider the following two scenarios (Travis, 1985, pp. 199–200; Travis, 1996, pp. 454–55; Searle, 1978, p. 212; Searle, 1980, pp. 224–225; Moravcsik, 1998, pp. 44–45): In the first, Smith has been dieting for the last eight weeks. He steps on the scale one morning, naked, before breakfast (but after having gone to the bathroom), and it registers 80 kg. A friend at work who wants to let Smith's co-workers in on his achievement can use (3) to say something true.

(3) Smith weighs 80 kg.

Notice that it doesn't matter at all that Smith is, at that time, dressed, wearing a heavy overcoat, and has just consumed an enormous lunch. In scenario two, Smith is exactly as in the first scenario. However, the speaker's circumstances (and purposes) have changed. At the time of this utterance of (3) (say, the same time as in the first scenario), Smith is about to enter an elevator with a capacity of no more than an extra 80 kg. An utterance of (3) in these circumstances could be both fatal and false. Note that what the scale registers when Smith is naked in the morning is in this context irrelevant.⁵

Simultaneous utterances of the same unambiguous, non-vague, non-elliptic sentence type with the same meaning, and whose referring terms have the same referents, are alleged to differ in truth-value; these philosophers conclude that (3) is context sensitive (see Travis, 1996, p. 455; Travis, 1997, p. 119; Searle, 1978, p. 219; Searle, 1980, p. 227). There is not supposed to be anything special about the linguistic items in (3); it is but one example radical contextualists advance; yet the intended scope is limitless (see Travis, 1996, p. 455; Searle, 1978, p. 219; Searle, 1980, p. 227).

Various brands of contextualism have been invoked to conclude that no sentence (or at least not many) semantically expresses a proposition or has (interpretive) truth conditions; to infer that our language harbors hidden or surprising indexicals, to resolve a debate between epistemic fallibilists and skeptics, to defend a brand of tolerable moral relativism, to unravel Sorites and Liar paradoxes, to protect a conservative view about psychological attitude attribution, to explain intuitions about quantifier domain under specification, and to provide a workable semantics for comparative and attributive adjectives.

With such lofty aspirations, it's too bad that contextualism fails. We will argue that the appeals to CSA typically invoked by contextualists are illegitimate. If they were legitimate, then certain other inferences ought to follow; since they do not, contextualism, at least in the cases in question, fails. A large part of our

discussion will be a defense and elaboration of what context shifting entails. We will conclude with two diagnoses about how contextualists got misled in the first place—in particular, we will argue that contextualists seduce themselves by vacillating between blurring the use/mention distinction, and introducing what Kaplan has called monsters into English. In effect, we will be defending a semantic invariantism about the controversial cases, whereby one is a semantic invariantist⁶ about an expression *e* just in case one rejects contextualism for *e*.

We begin by first drawing your attention to several general, natural and, we believe, obvious tests for linguistic context sensitivity. We show that the controversial cases fail each of the three tests. Throughout we use the expressions “know” and “red” as our examples of controversial cases. In the appendix we discuss comparative adjectives.

Test (i): Context Sensitive Expressions pass the Inter-Contextual Disquotational Test and permit Real Context Shifting Arguments

There are certain *live* tests a theorist can perform to determine whether an expression *e* is context sensitive. These tests are ‘live’ in this sense: they require the theorist to actually *use* *e* while performing the test. It requires the theorist to confront intuitions about her own language *in use*, not just about other people’s use of language (or her use in other contexts). To introduce the first live test, recall that a mark of a context sensitive *e* is that it can be used with different extensions (semantic values) in different contexts of utterance. It follows that we can use a context sensitive *e* in *this* context [i.e. the context of this paper] with an extension different than one it takes in another context. In effect, what we’re suggesting is that the theorist try to determine whether *e* is like that by actually using *e*.

Here’s what we have in mind. In order to use *e* we must first put it in a sentence *S* that we use. An expression *e* is context sensitive just in case there is some *S* containing *e* such that she (the theorist) can truthfully utter an instance of the following schema:

(ICD) There are (or can⁷ be) false utterances of ‘*S*’ even though *S*.

(Alternatively, run the test in reverse.⁸) Given this fact (see arguments for this claim below), to test for whether *e* is context sensitive or not, we see how our intuitions fall about utterances of legitimate instances of (ICD) for *e*. If it’s our intuition that no utterances of legitimate instances of (ICD) for *e* are true then one of two things follow. Either *e* is not context sensitive; or it is (and for some reason that requires explanation we cannot see that it is). If the latter, then we must reserve judgment on the context sensitivity of *e* until some such explanation is forthcoming.

So, for an illustration of an expression *e* that passes the (ICD) test consider us (i.e., Cappelen and Lepore) who are trying to determine whether ‘she’ is

context sensitive. To perform this task, we first choose a legitimate sentence *S* containing ‘she’, e.g., (4):⁹

(4) She is French.

(We’re actually pointing at a French woman right now.) We then proceed to assert (5a):

(5a) There is (or can be) a false utterance of ‘She is French’ even though she is French.

We then ask ourselves whether intuitively, what we’ve said by uttering (5a) (with the relevant intentions, demonstrations, etc.) is true. Is there a false utterance of (4) even though what we said when using (4) in our utterance of (5a) is true? Since the answer is ‘yes’—consider an utterance pointing at a woman who is not French—we’ve established that (4), and so ‘she,’ is context sensitive.

CSA proponents *must* accept ICD (and since this paper is about CSA, that’s what matters most to us). For, consider the *data* invoked in any CSA. The contextualist presents (simultaneous) utterances *u* and *u'* of an alleged context sensitive *S*, for example, ‘Lewis knows that penguins eat fish.’ The contextualist’s intuition is that *u* is true, while *u'* is false, even though the relevant facts about Lewis remain the same (he gains no new information, he does no additional reasoning, etc.). Whenever confronted with this data, we suggest you ask yourself: *OK, but is it or isn’t it the case that S?*—an obviously important question for someone who is *not* a contextualist, i.e., for someone who holds that instances of the following schema are true:

Utterances of ‘*S*’ are true just in case *S*.

That is, does Lewis know that penguins eat fish? If he does, *u'* secures a true ICD test utterance. If he doesn’t, *u* secures a true ICD test utterance (when the test is run in reverse). So clearly: *S is shown to be context sensitive by a CSA just in case S passes ICD.*

Here are additional illustrations involving uncontroversial context sensitive expressions:

(5b) There is a false utterance of ‘That’s nice’ even though that’s nice [said pointing at Al’s car].

Suppose Al’s car is nice. Then obviously, if there is an utterance where someone points at anything other than Al’s car that isn’t nice, then the test utterance expresses a truth.

(5c) There is a false utterance of ‘I’m hungry’ even though I am hungry.

Suppose I'm hungry. Then obviously, if there is an utterance of 'I'm hungry' where its speaker is someone other than me, who isn't hungry, then the test utterance expresses a truth.

- (5d) There is a false utterance of 'Tom is leaving now' even though Tom is leaving now.

Suppose Tom is leaving now. Then obviously, if there is an utterance of 'Tom is leaving now' made at times other than now, say, a few days into the future when Tom isn't leaving, then the test utterance expresses a truth.

As these stories decisively illustrate, there are uses of sentences with traditional indexical expressions that pass ICD with flying colors. That there are intuitively true utterances of (5a)–(5d) in perfectly ordinary circumstances suffices to convince us, not surprisingly, that *not* all utterances of the following bi-conditionals are true.

- 'She is French' is true just in case she is French.
 'That's nice' is true just in case that's nice.
 'I'm hungry' is true just in case I'm hungry
 'Tom is leaving now' is true just in case Tom is leaving now.

We will discuss below whether the alleged context sensitive expressions 'know', 'red', and 'weighs 80 kg' also pass ICD.

First, we introduce a closely related test we call the *RCSA-test*. We begin with some more stage setting. The context *in* which a CSA is told we'll call its *Storytelling Context*; the contexts *about* which a CSA is told we'll call a *Target Context*. In trying to invoke intuitions about context shifting we can devise either of two *sorts* of stories in a Storytelling Context: one in which the alleged context sensitive expression is *not* used but only mentioned in describing its uses in Target Contexts; or one in which it *is* used and also mentioned in describing its uses in Target Contexts. We'll call the former *Impoverished CSA* (ICSA) and the latter *Real CSA* (RCSA).

Contextualists invariably rely on ICSA to convince us that the relevant expressions are context sensitive. For example, to convince us that knowledge attributions are context sensitive, they appeal to intuitions we have about ICSA to provide evidence of context sensitivity, as in: consider two target contexts, one in which the topic of conversation is philosophical skepticism and one in which it is various issues about birds (nothing philosophical).¹⁰ Imagine an utterance of (1) in both.

- (1) Lewis knows that penguins eat fish.

Intuition is supposed to support the conclusion that the utterance of (1) in the first is false (because Lewis doesn't, for example, know how to rule out that he is

a brain in a vat), while the one in the second is true (since he's fairly knowledgeable about flightless water birds). So described this story is an ICSA since what an utterance of (1) expresses in the Storytelling Context is never asserted or denied; *we haven't told you whether or not Lewis knows that penguins eat fish*.

Notice that this ICSA differs from the stories we told about context shifting with regards to the sentences mentioned in (5a)–(5d); re-examine those stories and you'll see for yourself that each constitutes an RCSA for their mentioned sentences. (5a), for example, tells a story in which (4) is used in an assertion in the Storytelling Context.

In sum, for a story to be a RCSA for 'know' it can't just be *about* utterances of a sentence S containing 'know'; it must also affirm or deny what S expresses *in the Storytelling Context*. An ultimately misleading, but potentially illuminating way to couch this point is that an RCSA tells us whether or not the individual has the knowledge in question,¹¹ and an ICSA does not.

At this stage we want to register two points about RCSA and ICSA.

First, if an expression is genuinely context sensitive, we should be able to construct an RCSA for it, i.e., we should be able, in a Storytelling Context, to use it in an assertion and describe a Target Context in which it is used falsely (or *vice versa*); thus, convincingly establishing *bona fide* context shifting. It is only by telling such stories that we can establish that an expression passes the ICD test, and only such expressions, as already remarked, are context sensitive. If you don't already see this, note that according to a semantic invariantist each of the following bi-conditionals has only true utterances (ignoring tense):

'George knows that he has hands' is true just in case George knows that he has hands.

'Fire engines are red' is true just in case fire engines are red.

'Smith weighs 80 kg' is true just in case Smith weighs 80 kg.¹²

So, if contextualism is true, it must be established that these bi-conditionals have false utterances; otherwise, it's not at all clear what a claim of context sensitivity comes to. It follows¹³ that for any sentence S to be shown to be context sensitive it must be established both that it has (potentially) true utterances while denying what an actual use of S in the Storytelling Context expresses (or the other way around). That is, there must be an RCSA for S.

Second, if an ICSA is presented in defense of a contextualist claim, and if this ICSA cannot be transformed into an RCSA, we know that one of its premises is false: It's false either that we have (stable) intuitions about the truth value of the utterances referred to (they, so to speak, evaporate when placed in an alleged RCSA), or these intuitions will turn out to be semantically irrelevant.

Do the Controversial Cases pass ICD and admit of RCSA?

As with (5a)–(5d), the mentioned sentences in (6a)–(6b) are context sensitive only if there are true utterances of (6a)–(6b).¹⁴

- (6a) There are true utterances of ‘George knows that he has hands’ even though George doesn’t know he has hands.
- (6b) There are true utterances of ‘Fire engines are red’ even though fire engines aren’t red.

We deny, however, there are any such utterances. We believe that, even at this early stage of our argument, anyone who doesn’t already have theoretical prejudices will find it very hard to resist denying there are true utterances of (6a)–(6b). If any utterance of ‘George knows he has hands’ is true, George had better know he has hands; and if any utterance of ‘Fire engines are red’ is true, fire engines had better be red. These intuitions tell us that ‘know’, ‘red’, etc., fail the ICD test. Compounded with the intuitive obviousness of the ICD test for context sensitivity, this constitutes strong *prima facie* evidence that these expressions are *not* context sensitive.

Of course, it would be boring were the entire debate reduced to a collision of intuitions: we say all utterances of (6a)–(6b) are intuitively false; our opponents insist that they can hear some as true. How do we press forward? Well, since it’s supposed to be news that these expressions are context sensitive, anyone who thinks there can be true utterances of (6a)–(6b) needs to bolster her case: she could try to do so by bringing us to recognize some of these utterances as true, perhaps, by getting us to reflect further upon the sorts of data presented in CSA. Thinking about the cases presented by Cohen, De Rose, Lewis, and others might enable us to recognize there are true utterances of ‘George knows that he has hands’ even though George doesn’t know he has hands. Thinking about Travis’ and Searle’s examples involving ‘weighs 80 kg’ might enable us to recognize there are true utterances of ‘Smith weighs 80 kg’ even though Smith doesn’t weigh 80 kg. Etcetera.

It’s extremely telling, in this regard, that the stories presented in defense of contextualism *never* take the form of an RCSA; typically, we are given an ICSA. In what follows we will try to elicit just how hard it is, if possible at all, to devise RCSA for ‘know’, ‘red’, weighs 80 kg’, etc., that is, a story in which these words are both used (in the appropriate way) and mentioned. It’s crucial to keep in mind that our stories are contextualized, i.e., understood as uttered if you like. That’s important because if contextualism is true, then these RCSA will contain a context sensitive expression.

We call the first alleged RCSA *Known Rupert*:

Known Rupert

Right now, I’m doing philosophy and thinking about Rupert. Rupert, however, is not now doing philosophy. Instead, he’s home making tea. Rupert doesn’t

know he is thirty years old. For Rupert to know he is thirty years old, he *has to* rule out the possibility that he is a brain in a vat. Rupert, however, is unaware of (or not thinking about) this possibility.¹⁵ And so he's ignoring a possibility that *must* be ruled out in order for anyone to know anything at all. Still, when Rupert utters in the comfort of his home, 'I know I am thirty years old' what he says is true, because he's ignoring this possibility, even though this possibility has got to be considered in order for Rupert to know anything at all.

To see the point of the Known Rupert scenario remember that according to contextualism 'know' *is* context sensitive, and so, its semantic value is fixed in a context of utterance. When we *use* 'know' in the Storytelling Context (i.e. *this* context) to describe a Target Context, it takes on the semantic value it has in the Storytelling Context, and not the semantic value it would have had had it been used in the Target Context. After all, we are *not* in the Target Context; we're *in* the Storytelling Context using 'know' to describe the Target Context. More generally, when we *use* "a knows that p" in *this* context to describe a possible world, it is the standards of *this* context that determines whether an object in that possible world is correctly described by that utterance, i.e. whether *a* knows that p. That, by the way, is why Known Rupert contains the modal claim that Rupert is "ignoring a possibility that *must* be ruled out in order for anyone to know anything at all." Remember: *All of this follows directly from contextualism itself*. If 'know' is context sensitive, then Known Rupert should be true; we should have the intuition that it is true.¹⁶

Here's the same point applied to "red".

Red Rupert

In order to be red, an apple has to have red skin. That's a necessary condition for being a red apple. It is irrelevant, for instance, whether an apple is red on its inside. Here's an apple, call it Rupert; Rupert is red. On the inside, Rupert is white. Nonetheless, there are utterances of 'Rupert is red' that are false, not because Rupert's color changes, but because the speaker cares about what's inside Rupert rather whether it is red or not. This affects the truth-value of the utterance even though the color of the inside of the apple is completely irrelevant to whether Rupert is red.¹⁷

We assume you'll agree with us that these two Rupert stories do not provide clear and convincing intuitive support for contextualism. However, compare them with the following Now scenario (again, reading this passage as contextualized).

Now

Right now, Stephen is not wearing a hat. Yesterday he was wearing a hat. And when he then uttered 'I'm wearing a hat now' what he said then was true, even though he's clearly not wearing a hat now.

The Rupert scenarios are unconvincing to us, and compare quite unfavorably to the Now scenario. Unlike the Rupert scenarios, we take Now to be an RCSA that provides evidence for the view that ‘now’ is context sensitive. Yet there is nothing cagey about the Rupert stories; they parallel exactly the Now one. If you don’t like them, or think that they are prejudicially slanted, try devising one of your own.

Test (ii): An Expression is Context Sensitive only if it Blocks Inter-Contextual Disquotational Indirect Reports (IDI-Test)

Here’s what we mean by Inter-Contextual Disquotational Indirect Reports:

Take an utterance *u* of a sentence *S* by speaker *A* in context *C*. An inter-contextual disquotational indirect report of *u* is an utterance *u'* in a context *C'* (where *C' ≠ C*) of “*A* said that *S*.”¹⁸

An expression in *S*, e.g., *e*, blocks an inter-contextual disquotational indirect report just in case its occurrence in the complement clause of the indirect report typically would render the indirect report false. If this sounds confusing, it shouldn’t. Take an obviously context sensitive expression, e.g., the first person pronoun “*I*”. Consider an utterance of the sentence “*I* went to Paris” by Rupert. If Lepore tries to report what Rupert said with “Rupert said that *I* went to Paris,” his report is false because the expression “*I*” fails to pick out what “*I*” picked out in the original utterance. The presence of “*I*” in the disquotational report figures prominently in an explanation of why the report is false. In general, the presence of “*I*” in the subject position in the original utterance cannot be preserved in an indirect report, i.e., you can’t report disquotationally (there are of course obvious and easily explainable exceptions to this generalization: for example, self-reporting.)

It’s (almost) a matter of definition that context sensitive expressions tend to block inter-contextual disquotational indirect reports. By definition, to restate the obvious, for *e* to be context sensitive is for *e* to shift semantic value from one context of utterance to another. So, if *e* is context sensitive and Rupert uses *e* in context *C*, and Lepore uses it in context *C'*, then it will be just an *accident* if their distinct uses of *e* end up with the same semantic value. In particular, if Lepore finds himself in a context other than Rupert’s and wants to utter a sentence that matches the content of Rupert’s utterance of a sentence containing *e*, he can’t, typically,¹⁹ use *e*, i.e., he can’t report Rupert’s utterance disquotationally.

It is rather obvious, we think, that “red” and “know” do *not* block inter-contextual disquotational indirect reports. If *B* utters “*A* is wearing a red hat,” there’s no *prima facie* case against using “red” in the indirect report, as in “*B* said that *a* is wearing a red hat.” It’s not at all like re-using “*I*”, or “that.” Similarly for “know”; if *B* utters “*A* knows when the train leaves,” there is no *prima facie*

case against using “know” in a disquotational indirect report, as in “B said that A knows when the train leaves.”

Test (iii) Collective Description Test for Context Sensitivity

If a verb phrase *v* is context sensitive (i.e., if it changes its semantic value from one context of use to another), then on the basis of merely knowing that there are two contexts of utterance in which ‘A *v*-s’ and ‘B *v*-s’ are true respectively, we *cannot* automatically infer that there is a context in which ‘*v*’ can be used to describe what A and B have both done. In short, from there being contexts of utterance in which ‘A *v*-s’ and ‘B *v*-s’ are true it doesn’t *follow* that there is a true utterance of ‘A and B both *v*.’ This is because the semantic value of ‘*v*’ in the previous sentence is determined in one context, and we have no guarantee that that semantic value, whatever it is, ‘captures’ (whatever that means) the semantic values of ‘*v*’ in those contexts of utterance where they were used alone.

On the other hand, if for a range of true utterances of the form ‘A *v*-s’ and ‘B *v*-s’ we obviously *can* describe what they all have in common by *using* ‘*v*’ (i.e., by using ‘A and B *v*’), then that’s evidence in favor of the view that ‘*v*’ in these different utterances has the same semantic content, and hence, is not context sensitive.²⁰ A parallel point extends to singular terms.

If a singular term *N* is context *insensitive* and there’s a range of true utterances of the form ‘N is F’ and ‘N is G’, then we, for example, in *this* context, can truly utter ‘N is F and G.’ Similarly, if *N* is context sensitive, we shouldn’t be able to do this. As an illustration consider the context sensitive ‘yesterday’: Suppose we know there are two contexts in which ‘Yesterday John left’ and ‘Yesterday Bill left’ are true respectively (though we don’t know the time of these contexts). It doesn’t follow that there is a context in which ‘Yesterday John and Bill left’ is true.

Let’s apply this test to the controversial cases. We start with ‘know.’ Consider these three utterances, and suppose all of them are true:

A knows that penguins eat fish.

B knows that penguins eat fish.

C knows that penguins eat fish.

Imagine these utterances are accompanied by typical contextualist stories, stories that are supposed to lead to your seeing that the context of utterance affects the relevant epistemic standards, hence, affects the semantic value of ‘know.’ Suppose, for example, A is a five year old, and it’s his (non-philosophical) mother who truly attributes knowledge to him. Suppose B is biology major, and it’s his (again, non-philosophical) biology professor who truly attributes the knowledge to her, and that C is a philosophy graduate and it’s his philosophy professor who truly attributes the knowledge to him as a way of saying that he

has succeeded in responding to skepticism. Here's a procedure for evaluating claims about contextual variation:

If, having heard the contextualist stories, you still think it follows that A, B and C all know that penguins eat fish, then you have evidence 'know' isn't context sensitive.

Our strong intuitions are that A, B and C know that penguins eat fish. So, contextualist stories provide no evidence for context sensitivity. Moreover, we have direct evidence in favor of the view that 'know' is context *insensitive*.²¹

For a second illustration consider utterances of:

That's a red car.
That's a red snake.
That's a red house.
That's red hair.

Again imagine these utterances accompanied by contextualist stories that lead you to think the claims made by these utterances are true. Now, ask yourself: Is there something the car, the snake, the house and the hair have in common? Yes. They are all red. That's obviously true, isn't it? If we can use "red" to describe what the car, the snake, the house and the hair have in common it means that "red" doesn't pass the Collective Description Test for context sensitivity.

Possible reply: Contextual Salience Absorption

The sort of reply we hear most often (to all of these objections) goes something like this: schematically, the situation a theorist (or a couple of theorists, as in our case) is in is this: You are in a context 5stC. (This is the name of the context we wrote this paper in, a café on 5th Street between Avenues A and B in Manhattan.) In 5stC, you imagine other contexts, e.g., C1 and C2 in which someone utters, say, (7).

(7) *A is red.*

You then ask yourself what is said by these two utterances of (7). You do that by vividly imagining C1 and C2 in 5stC. You imaginatively place yourself first in C1 and then in C2. To imaginatively place yourself in C1 and C2 triggers what might be called *contextual salience absorption*. When in 5stC, you imagine yourself in C1, and then what is salient in C1 becomes salient to you. When you imagine yourself in C2, what's salient in C2 then becomes salient. So, what's contextually salient in 5stC varies contingent upon what you happen to be thinking about in 5stC. So far so good. But why is this a response to our three objections?

In each of the three cases, the theorist is asked to vividly imagine a target context from a storytelling context. If imagining a context, say, C1, from the perspective of another context, say, 5stC, renders what is salient in C1 salient in 5stC, then that undermines the idea that separate contextual standards are in play. The three tests for context sensitivity (in different ways) all assume that the contextually salient standards are clearly distinct. The three tests play on the distinction between, e.g., the used “know” and the mentioned “know.” To say that there’s contextual salience absorption is to say that the standards are blurred. That would make the tests unreliable. To put it pointedly: as soon as you think about a context C, the standards of C affect the standards of the context you’re in (the context in which you do your thinking). Just thinking about C changes the context you are in so that there is no longer a clear distinction in contextual standards between the storytelling context and the target context. If this is so, our tests are faulty: they all presuppose that there is a rather sharp distinction between the storytelling standards and the target context standards.

We have left this objection rather vague in part because we’ve never seen it worked out in detail and there are many ways to do that. But, we think, no matter how it’s elaborated this response will fail miserably for two reasons.

- (a) The reply gets the dialectic deeply wrong and leaves the contextualist with no arguments against Semantic Invariantism.
- (b) It triggers what we think ultimately leads to an internal inconsistency in various contextualist positions.

A. An Appeal to Absorption Leaves the Contextualist Dialectically Impotent

Let’s review the dialectical situation we are in: a semantic invariantist about an expression *e* is someone who thinks there’s no contextual variation in the semantic content of utterances of sentences containing *e* (other than that contributed by expressions other than *e*). The contribution of *e* to the proposition semantically expressed by utterances of sentences containing *e* is the same in every context of utterance. One way to capture this view (not the only way, but it captures a common core of many versions of Semantic Invariantism) is this:

If *S* is a sentence containing *e* and no other indexical components, then an utterance of either (I) or (II) should be true:

- (I) “*S*” is true just in case *S*.
- (II) “*S*” expresses the proposition that *S*.

Context shifting arguments are supposed to convince a semantic invariantist that her position is flawed, i.e., they are supposed to provide evidence *against* semantic invariantism. The evidence is supposed to take the form of a counter-example: intuitive evidence that there is at least one utterance *u* of *S* that

semantically expresses a different proposition, has different interpretive truth conditions (maybe even has a different truth-value with respect to the same circumstance of evaluation) as an utterance of *S* in Storytelling Contexts. To endorse contextual salience absorption is in effect to grant that no such counter example is forthcoming. Whenever in the Storytelling Context we think about or describe a Target Context, our intuitions about the e-standards of the Storytelling Contexts and Target Contexts are unified or blurred. To get a clear counter example of this kind the contextualist presupposes clearly distinct standards. In sum: appeal to contextual salience absorption is not a defense of contextualism; rather, it's a concession of defeat.

B. Contextual Absorption Renders Contextualism Internally Inconsistent

A corollary of point (A) is this: if we assume for the sake of argument that a sentence *S* is context sensitive, and if the context sensitive component of *S* absorbs, then there's overwhelming evidence in favor of the truth of (D1) and (D2):

(D1) Every utterance of "A is red" true iff A is red.

(D2) Every utterance of "A is red" semantically expresses the proposition that A is red.

Remember, if "A is red" is absorbent, then every utterance of "A is red" can be characterized truly by an utterance of (D1.1) and (D2.1):

(D1.1) "A is red" is true iff A is red.

(D2.1) "A is red" expresses the proposition that A is red.

Every utterance of (D1.1) and (D2.1) is true because, if "A is red" is absorbent, then the truth conditions for "A is red" are the same in the Storytelling Context and the Target Context. The truth of (D1.1) and (D2.1) provides overwhelming evidence in favor of the generalizations (D1) and (D2). But (D1) and (D2) are the denial of contextualism. In other words, a contextualist about *S* who also claims that *S* is absorbent has an incoherent view.²²

Diagnosis: Why Impoverished Context Shifting Arguments are Seductive

If we are right, then "know" and "red" fail three obvious tests for context sensitivity. We think the same applies to every other example mentioned at the beginning of this paper. It even applies to comparative adjectives (see appendix). So what's going on? How do philosophers convince themselves that obviously non-context sensitive expressions are context sensitive?

These positions are typically justified by an appeal to *Impoverished Context Sensitive Arguments*. That raises the question: if you have an Impoverished

Context Shifting Argument for an expression *e*, but *e* does not pass other tests for context sensitivity (including the three tests discussed here), what's going on?

First of all, you have strong evidence that whatever intuitions were triggered by the ICSAs concern non-semantic content. They are intuitions that reveal what the utterance might succeed in communicating, but not its semantic content. But there's more than that going on. We have two tentative diagnoses of why contextualists are so easily seduced by their ICSAs.

First Diagnosis: Treating 'in context C' as a Monster

Contextualists talk quite freely about knowing that *p* *in context C* and not knowing that *p* *in context C'*, or being rich *in C* but not being rich *in C'*, etc., as if the semantic role of 'in context C' were to map *S* (the sentence to which it attaches) from a context of utterance onto a distinct context of utterance. Unwittingly, we presume, in this regard they are treating 'in context C' as what David Kaplan has called a monster. A monster is an operator...

which when prefixed to a *sentence* yields a truth if and only if in some contexts the contained sentence (not the content expressed by it) expresses a content that is true in the circumstance of that context. (Kaplan, 1989, p. 510)

If 'in context C' were a monster, it could transport us from our current context of utterance to imagined ones. It would allow us to *use* an expression in this context, within the scope of the monster, to say what that expression would say if it were uttered in another context. This, so it seems to us, is what often transpires when contextualists try to devise (fake) RCSA. And, of course, if English had monsters and 'in context C' was one of them, then most of our objections would be undermined. For example, you would be able to use 'know' or 'red' in a Storytelling Context *C* when describing a Target Context, but they wouldn't take on the semantic values they have in *C*. If there were monsters and contextualists could use them, the effect would be this: *the theorist could keep a sharp distinction between the standards of the storytelling context and the target context, and could describe what goes on in the target context either by using, say, "know", outside of a monster or inside a monster. With this device in hand, she could say one thing within the monster and another outside it. Playing on this discrepancy she get herself to believe she has something like a RCSA rather than the obviously unsatisfying ICSA.*

Contextualists certainly write as if they can use 'know' or 'red' in this way. But you could succeed with these uses only if there were monsters, and, as Kaplan has pointed out, 'in context C' is *not* a monster in English. He writes:

Let us try it:

(8) In some contexts it is true that I am not tired now.

for (8) to be true in the present context it suffices that some agent of some context not be tired at the time of that context. (8), so interpreted, has nothing to do with me or the present moment. But this violates principle 2!²³ Principle 2 can also be expressed in a more theory-laden way by saying that indexicals always take primary scope. If this is true—and it is—then no operator can control the character of the indexical within its scope, because they will simply leap out of its scope to the front of the operator. I am not saying we could not construct a language with such operators, just that English is not one. And such operators could not be added to it (Kaplan, 1989, p. 510).²⁴

Indirectly, we have been championing Kaplan's view. Our first diagnosis of what's gone wrong with contextualism is that in convincing themselves that the expressions in question are context sensitive, the locution 'in context C' winds up being treated as a monster. A case in point is Graff's discussion of color words:

Suppose I want you to hand me a certain book. If the book in question is colored a very light grayish-blue, and it's sitting amongst a bunch of other books all of which are colored a very light grayish-red, I may say, 'Hand me the blue one.' If, on the other hand, the book I want is sitting with a bunch of richly-colored cobalt blue books, I may say, 'Hand me the gray one.' I take it that it would be true to say in the first case that the book I wanted was blue, and in the second case that the book I wanted was gray (Graff, 2000, p. 56).

Focus carefully on the last sentence of her story: 'I take it that it would be true to say in the first case that the book I wanted was blue, and in the second case that the book I wanted was gray.' Call its tokening *u* and the context of *u* the Storytelling Context. The color words 'blue' and 'gray' are *used* there in an indirect quote that occurs in the Storytelling Context.

Context sensitive expressions in indirect quotes receive the semantic value of the context of the indirect quote itself. If we say, in the context of this paper, that Jason said our view is crazy 'our' refers to Cappelen and Lepore, not to Jason and company. Indirect quotation is no monster. In other words, Graff's use of 'blue' and 'gray' get their semantic values from the Storytelling Context and not from Target Contexts. But then *u* makes no sense. Graff's point is that the semantic values of 'blue' and 'gray' shift between what she calls *the first* and *the second case* (we call these the first and second Target Contexts). That wouldn't work were both given their interpretations in the Storytelling Context. She wants 'in the first case' to create a context such that what occurs after it is interpreted as uttered in the first Target Context. She wants what occurs after 'in the second case' to be interpreted as uttered in the second Target Context. On this interpretation, she is treating 'in the first (second etc) case' as monsters.²⁵ But again, there are *no* monsters in English. When someone says to you, 'In case

one, I'm not happy', there is no context in which his use of 'I' picks out you—which it should be able to do were 'in case one' a monster.

For a second illustration, consider a passage from Soames (1998), where he's endorsing the view that 'looks green' is context sensitive. Here's how he describes the context sensitivity:

In this model the rule governing contextual adjustment of the predicate *looks green* should be roughly as follows:

If in a context C something x is explicitly judged to look green, then the extension of *looks green* in C includes everything perceptually indistinguishable in color from x, as well as everything that looks greener than x. If in C something y is judged not to look green, then the anti-extension of *looks green* in C includes everything perceptually indistinguishable in color from y, as well as everything that looks less green than y. (Soames, 1998, pp. 211–212)

This entire statement is uttered in the context of Soames' book. Call it the SB-Context (for Soames' Book Context). Focus on the first conditional and the first use of 'to look green' in that conditional. Call this utterance of 'to look green' u (so u is preceded by an utterance of 'If in a context C something x is explicitly judged...'). Now ask yourself: What's the semantic value of u? The expression 'to look green' is, according to Soames, context sensitive, so its semantic value is determined by its context of utterance. The context of utterance in this case is the SB-context. So u should get its semantic value from SB. But that's clearly not what Soames intends. For then he would end up talking only about contexts in which someone uttered (or made a judgment) about things that fall under 'looks green' according to the standards of the SB-Context. What Soames wants is for u to take its semantic value from the imagined C. He wants u to have the semantic value it *would have had* if it had been uttered in this imagined context. That's why he puts 'If in a context C something x is explicitly judged...' in front of u. In other words, he is treating 'in a context C' as a monster.²⁶ But, again, there are no monsters in English; and surely, it should be obvious that 'in a context C' is not one.²⁷

For a third illustration, consider this story from Cohen:

Mary and John are at the LA airport contemplating taking a certain flight to New York. They want to know whether the flight has a layover in Chicago. They overhear someone ask if anyone knows whether the flight makes any stops. A passenger Smith replies 'I do. I just looked at my flight itinerary and there is a stop in Chicago.' It turns out that Mary and John have a very important business contact that they have to make at the Chicago airport. Mary says: 'How reliable is that itinerary, anyway. It could contain a misprint. They could have changed the schedule since it was printed, etc.' Mary and John agree that Smith doesn't really know that the plane will stop in Chicago on the basis of the itinerary....(Cohen, 1999)

Cohen presents his thought experiment in the Storytelling Context. Consider the utterance *u* of ‘They overhear someone ask if anyone knows whether the flight makes any stops.’ Since *u* is uttered in the Storytelling Context, if contextualism is true, ‘know’ in *u* takes on a semantic value in that context. Suppose the proposition expressed by *u* in the Storytelling Context is *p*. Now consider the direct quote of Smith that directly follows *u* (where Smith is quoted as saying ‘I do’). Call the utterance of this direct quote *u'*. Cohen clearly thinks *u'* indicates Smith agrees with the semantic content of *u*. But it’s a mystery how it could. *u* is uttered in the Storytelling Context, the context of a philosophy paper, and, again assuming contextualism, the fact that *p* is expressed by *u* is due to peculiar features of that context. Smith’s utterance *u'*, however, is made in the Target Context, a non-philosophical context (i.e., a context with low epistemic standards). The epistemic standards for ‘know’ are fixed by the kind of considerations Cohen and other philosophers are concerned with. Since Smith is not in that sort of context, how could he end up agreeing with *p*? It’s all very peculiar. What’s going on? One possibility is that Cohen wants us to read *u* as if it was uttered in the Target Context. To do so would be to treat the act of storytelling as tacitly introducing a monster.²⁸

Our own view about ‘in a context *c*’ is that it has no impact on the semantic value of the controversial cases. If an object is round in room 300, it’s round. If someone knows in context *c* what penguins eat, then he knows what penguins eat.

Second Diagnosis: How to Avoid Monsters: Direct Quotation

Aren’t we just being pedantic? Couldn’t a contextualist easily avoid these troubles just by being more careful? In particular, couldn’t they avoid monsters simply by using direct quotation instead? That’s probably what they think and it’s probably why they allow themselves to be so sloppy. Here’s David Kaplan on the connection between monsters and direct quotation:

There is a way to control an indexical, to keep it from taking primary scope, and even to refer it to another context. Use quotation marks. If we mention the indexical rather than use it, we can, of course, operate directly on it. Carnap once pointed out to me how important the difference between direct and indirect quotation is in

Otto said ‘I am a fool’.

Otto said that I am a fool.

Does direct quotation provide an easy way out for the monstrous contextualist? No. If in order to avoid monsters a contextualist resorts to direct quotation, then she would be left with nothing more than intuitions about an ICSA to bolster her case. But as we have been arguing, those intuitions alone are insufficient to establish contextualism unless a corresponding intuitive RCSA

can be devised. But an RCSA requires *using* the disputed expression in an assertion (or a denial). It is tempting to do that in a monstrous way (in which case the contextualist would have her RCSA), but there are no monsters! The way to avoid monsters is to resort to direct quotation, but then the contextualist is back to an ICSA.

Conclusion

There can't be metaphysical, epistemological, or moral arguments for an expression being context sensitive. If an English expression *e* is context sensitive, then that's a fact about English and it can only be established by whatever procedures semanticists employ to establish such facts. A cornerstone of semantic methodology is the appeal to intuitions. These intuitions tend to be about language use. The assumption is that as competent speakers we at some level know the semantics for *e* and in order to retrieve semantics facts about *e* all we need to do is remind ourselves of how it is used in various contexts of utterance. Here's a way to think about the arguments in this paper: intuitions about imagined contexts of utterance are determined by two factors: the linguistic facts about those imagined contexts (i.e., the sentences uttered in them) and the non-linguistic facts about those imagined contexts. How you describe the non-linguistic facts determines your intuitions about the content (truth value and truth conditions) of language used in those imagined contexts. The common thread that runs throughout our arguments is that contextualists tend to forget that semantics itself is done in context. The fact that our imaginings of other contexts of utterance are done in a context constrains how we can describe these imagined contexts and that again determines what intuitions we have. If we are right, then intuitions that at first blush might seem to provide evidence for context sensitivity tend either to evaporate or weaken when properly contextualized. Radical contextualists tend to be oblivious as to how their own arguments are contextualized. In sum: contextualists don't take context seriously.

Appendix on Comparative Adjectives

Contextualism about comparative adjectives is commonplace, and again the typical defense is an appeal to context shifting (e.g., Parsons 1972, Kamp 1975, Ludlow 1989). Consider two contexts, one in which the topic is the heights of Saudi Arabians, the other the heights of NBA basketball players. In such circumstances, it is argued that utterances of (9) disagree in truth-value.

(9) Osama Bin Laden is tall.

This is not so because Osama has grown or shrunk, or because others have grown or shrunk, but simply because the *standard of comparison* has shifted from one context to the other.

Comparative adjectives are context sensitive just in case they pass our three tests. In this appendix, we first argue that they *don't* pass these tests. We then consider two possible reactions to this failure. Throughout we use 'rich', 'expensive' and 'tall' as our examples of comparative adjectives.

ICS/RCSA-Tests for 'rich'

Can an RCSA be developed for "rich"? We don't think so. Consider this little story:

Rich Rupert

Rupert isn't rich. Anyone who is rich must make more money than 92% of Americans.²⁹ That's necessary for being rich. It is impossible to be rich without satisfying this condition. Rupert doesn't. His income is average. However, one could still utter 'Rupert is rich' and express a true proposition, but not because Rupert makes more money relative to other Americans, but rather because in this possible context of utterance, some other comparison class has been rendered salient. This other comparison class, however, is of course irrelevant to whether or not Rupert is rich; again, a person is rich just in case he makes more than 92% of Americans. Still, somehow or other, the salience of this other comparison class (the comparison class that's irrelevant to whether or not he is rich) makes a difference.

As in the other cases, the story seems peculiar, to say the least. Why would the salience of a comparison class that's irrelevant to Rupert's richness suddenly render an utterance of "Rupert is rich" true, when Rupert isn't rich? We doubt that anyone could (or would seek to) find solid support for the context sensitivity of "rich" in this story.

Collective-Description-Test for "Expensive" and "Tall"

Here are two illustrations of the Collective Description test using a comparative adjective. Suppose following Harriet around all day, observing her shopping practices, we notice that in each store when she's done shopping she always tells us 'I bought something expensive'—whether she's buying shoes, jewelry, a car or whatever. If contextualism about 'expensive' were true, it would *not* follow that everything she bought is expensive (or that *they* are expensive (where "they" refers to the things she bought)). The previous sentence should mean that she buys things that are expensive according to the standard, norm or comparison class invoked in this paper, since that's at most one context, her collective use shouldn't come out true. But that result is deeply counter-intuitive. Isn't it true that everything she bought was expensive?

Or consider a bunch of tall things, for example, the Empire State Building, Michael Jordan, and Mount Everest. Do these tall things have something in common? Yes. They are all tall. Alternatively imagine a range of true utterances:

Michael Jordan is tall.

Mount Everest is tall.

Empire State Building is tall.

It seems obvious that we can say, knowing that these utterances are true, that *they are tall*. Hence “tall” does not pass the collective description test for context sensitivity.³⁰

Inter-Contextual Disquotational Indirect Report Test for “Tall” (IDI-Test)

Suppose you hear a snippet of a conversation, you might just hear B utter “Rupert is tall.” Suppose you know very little about the rest of the conversation. You’re then asked what B said. There’s an answer, isn’t there? You’re perfectly justified in responding that B said that Rupert is tall. You’re justified in saying this because it’s true. So “tall” doesn’t pass the IDI-test.

Possible Responses from those Committed to the Idea that Comparative Adjectives must be Context Sensitive

Some philosophers are deeply devoted to the idea that comparative adjectives are context sensitive. They are so devoted for a range of reasons. Limitation of space prevents a full-scale discussion of these additional reasons here (but see Chapter 10 of our *Context and Sensitivity* for an extended discussion). However, whatever your reasons might be for thinking that comparative adjectives are context sensitive, the fact (if it is a fact) that they fail these three tests forces a contextualist about comparative adjectives to do either of two things:

- a) Show that we’re wrong in claiming that comparative adjectives don’t pass these three tests; or
- b) Show that an expression can be context sensitive even though it doesn’t pass the three tests.

We don’t see how to execute either strategy, but we’re open-minded.

Notes

1. Earlier drafts of this paper were presented to various audiences. We would like to thank them all. Special thanks to Ann Bezuidenhout, Kent Bach, Emma Borg, Manuel García-Carpintero, Dan Blair, David Braun, David Chalmers, Keith DeRose, Ray Elugardo, Jerry Fodor, Michael Glanzberg, Gil Harman,

- John Hawthorne, Kent Johnson, David Kaplan, Sarah-Jane Leslie, Peter Ludlow, Kirk Ludwig, Europa Malynicz, Terry Parsons, Jim Pryor, Mark Sainsbury, Stephen Schiffer, Rob Stainton, Jason Stanley, Brian Weatherson and Doug Winblad.
2. Cohen writes, ‘...when I say ‘contexts’, I mean ‘contexts of ascription’. So the truth-value of a sentence containing the knowledge predicate can vary depending on things like the purposes, intentions, expectations, presuppositions, etc., of the speakers who utter these sentences’ (1999, p. 57). De Rose writes, ‘the truth conditions of sentences of the form ‘S knows that p’ or ‘S does not know that p’ vary in certain ways according to the context in which they are uttered’ (1992, p. 914; cf., also, De Rose, Sec. 8, 1995; Casteneda 1980).
 3. Dreier writes, ‘speaker relativism is the theory that the content of (what is expressed by) a sentence containing a moral term varies with (is a function of) the context in which it is used’ (1990, p. 6).
 4. Context shifting does not require a change in truth-value. Consider the sentence ‘Every bachelor is unmarried’. According to contextualists its distinct utterances might range over distinct domains but the truth-value of each will be the same, i.e., true.
 5. Bezuidenhout writes, ‘We’re at a county fair picking through a barrel of assorted apples. My son says ‘Here’s a red one’, and what he says is true if the apple is indeed red. But what counts as being red in this context? For apples, being red generally means having a red skin, which is different from what we normally mean by calling a watermelon, or a leaf, or a star, or hair, red. But even when it is an apple that is in question, other understandings of what it is to call it ‘red’ are possible, given suitable circumstances. For instance, suppose now that we’re sorting through a barrel of apples to find those that have been afflicted with a horrible fungal disease. This fungus grows out from the core and stains the flesh of the apple red. My son slices each apple open and puts the good ones in a cooking pot. The bad one he hands to me. Cutting open an apple he remarks: ‘Here’s a red one’. What he says is true if the apple has red flesh, even if it also happens to be a Granny Smith apple’ (2002, p. 107).
 6. Unger (1984) introduced the term to talk about someone who denies that relevant contextual factors can affect the truth conditions of knowledge attributions, but we’ll use it for someone who denies contextualism in general.
 7. In our formulation of (ICD) we appeal to possible utterances. We do that because it is something our opponents, i.e., contextualists, do all the time. However, if you’re worried about quantification over possible utterances, (or worried it will be difficult to specify the relevant domain of possible utterances) run the test on actual (past, present or future) utterances.
 8. I.e., it can take the form of an utterance of ‘There is at least one true utterance of ‘S’ even though it is not the case that S.’
 9. Of course, care must be practiced in choosing S; e.g., S shouldn’t contain any other candidate context sensitive expression other than e, or, if it does, restrict the domain of ‘There are utterances’ so that additional context sensitive expressions must lock on to the same semantic values as in the test utterance.
 10. As Keith DeRose pointed out to us (personal communication), he prefers to use examples in which the high standard context is non-philosophical, e.g., a context in which the participants care very much about the evidence for the knowledge

claim. The exact points we make here extend to such examples by just changing the descriptions of the contexts appropriately. As far as we can tell, nothing at all hinges on what kinds of high/low standard examples are chosen. We let the 'high' standard cases be philosophical just because that is how Lewis proceeded.

11. It's misleading because here we are using 'knowledge' in the context of *this* paper, and we're not requiring that others use it to mean what it means in this paper (assuming that contextualism is true).
12. If you like to theorize about propositions semantically expressed, rather than interpretive truth conditions, CSA are supposed to show that these views are wrong:

(Every utterance of) 'George knows he has hands' expresses the proposition that *George knows he has hands*.

(Every utterance of) 'Fire engines are red' expresses the proposition that *fire engines are red*.

(Every utterance of) 'Smith weighs 80 kg' expresses the proposition that *Smith weighs 80 kg*.

(In what follows we focus on truth conditions since if two utterances differ in truth conditions, they express distinct propositions.)

13. Not quite immediately: the contextualist also needs certain assumptions about sentences and utterances. We won't go into that here.
14. To repeat why: Suppose 'George knows he has hands' is context sensitive, i.e., that the proposition expressed by (and the truth conditions of) 'George knows that he has hands' varies across contexts of utterance. If so, this sentence in *this* context [i.e., the context of this paper] expresses a certain proposition, and has certain truth condition. This proposition and these truth conditions needn't be the same as those of its other utterances. In other words, it's a trivial implication of the assumption that 'George knows that he has hands' is context sensitive that it has (at least potential) utterances that are not true just in case George knows he has hands. At least one of these is true even though George doesn't know he has hands, i.e., some utterance of (7a) [in the context of this paper] is true.
15. We've been assured by contextualist friends that no contextualist would require Rupert to occurrently rule out the possibility of being a brain in the vat, or to actively check it off in any manner. All that would be required is that Rupert is disposed to handle the possibility properly, perhaps by being capable of eliminating it on the basis of his evidence. We assume nothing in our thought experiment turns on this distinction.
16. The point can be strengthened: Suppose someone reads Known Rupert and says: 'Yeah, I can still hear that as true.' So far we have nothing but an incredulous stare to reply with. But we could catch our breath and go further. We can ask this person: why do you think Rupert's utterance of 'I know how old I am' is not true in this context? The reply, we suppose, is going to be something like: 'Because Rupert doesn't know he's not a brain in a vat.' But now ask: what is the meaning of 'know' in that reply? Is it what it means in this context or in the context of utterance? It can't be either, so our opponent is in a bind. What this shows is that in order for a CSA to be effective, it must be motivated by

descriptions that are context insensitive. (We'll see this point again when we discuss color words and their alleged context sensitivity below: color terms cannot be used to describe the Target Context in a CSA that aims to establish that such words are context-sensitive, because if they are, those words would take on the meaning they have in the Storytelling Context.)

17. *Skinny Rupert*

Rupert has been dieting for the last eight weeks. Rupert now weighs 80 kg! In order to weigh 80 kg a person must weigh 80 kg on an accurate scale, naked, before breakfast, in the morning. What he weighs with his clothes on at lunch is irrelevant. It has no bearing on whether or not Rupert weighs 80 kg. Nonetheless, there are utterances of 'Rupert weighs 80 kg' that are false, not because Rupert weighs more naked before breakfast in the morning, but because the speaker cares about what a scale would show when he steps on fully clothed after lunch. Suppose, for example, Rupert is about to get on an elevator with a capacity of no more than an extra 80 kg. If someone were to utter 'Rupert weighs 80 kg' her utterance would be false, even though he weighs 80 kg. The utterance would be false, not because Rupert's weight has changed, but because the speaker is concerned with something other than what Rupert weighs, for example what a scale would register were he to step on it fully clothed.

18. Of course, uttered as a report of *u* (if you want to make that explicit in the report just add "by uttering *u*" after "*S*").
19. Of course, he might accidentally find himself in a context that happens to assign the same semantic value to *E* as Rupert's context. But that would be an accident.
20. The argument can be summarized as follows: If '*v*' is a context sensitive term, then its semantic value changes from one utterance to another. So '*a v-s*' and '*b v-s*' attribute different properties to *a* and *b*. But it doesn't follow that '*v*' can be used to describe what *a* and *b* share. Maybe by chance someone might be able to use '*v*' in some context to refer to a property they all share, but that would be a coincidence. In other words: *e* is context sensitive just in case there's no guarantee of collective usage. Suppose there was a guarantee of collective usage, then a use of '*v*' in one context would 'denote' (have as its semantic value) what all other utterances of '*v*' denote and we would be guaranteed collective descriptions.
21. This test is closely related to ICD and RCSA. (By some methods of individuation, they might even be making the very same point.) If, for example, 'know' doesn't pass ICD, then every true utterance of the form '*A* knows that *p*' is true just in case *A* knows that *p*. Hence, any utterances of 'Tom knows that *p*' and 'Sally knows that *p*' are true just in case Tom and Sally knows that *p*.
22. There's a third worry for a contextualist who endorses contextual salience absorption: she has to explain why that does not undermine the possibility of RCSAs for "you" or "that" or other classical indexicals.
23. I.e., the thesis that indexicals pick out their referents directly from the context of utterance, without mediation. This means that the value of an indexical is fixed by the context of its utterance, and cannot be changed by the logical operators in whose scope it may occur.
24. Lewis, though he doesn't use the terminology, is acknowledging the same point when he writes: "To be sure, we could speak a language in which "As for you, I

am hungry” is true iff “I am hungry” is true when the role of speaker is shifted from me to you—in other words, iff you are hungry. We could but we don’t. For English, the speaker is not a shiftable feature of context” (Lewis, 1980, pp. 27–28).

25. Alternatively, she treats ‘what was said’ as a monster. The same point about monsters extends to De Rose, when he writes: ‘In Bank Case B...when, in the face of my wife’s doubt, I admit that I don’t know that the bank will be open on Saturday, I don’t contradict an earlier claim to know that I might have made before the doubt was raised and before the issue was so important because, in an important sense, I don’t mean the same thing by ‘know’ as I meant in the earlier claim’ (De Rose, 1992, p. 921).
26. Note also the following passage from Soames: ‘What, then, is going on? The answer, in our analysis, is that there is no contradiction here at all. According to the analysis, when x_i is initially characterized as looking green, this is done with respect to a certain set of standards, S . Later, when it is characterized as not looking green, this is done with respect to a new set of standards, S^* . But there is no contradiction in the observation that something may look green with respect to one set of standards and not look green with respect to a different set of standards’ (1998, p. 313). Note the occurrence of ‘looking green’ in ‘when x is characterized as looking green’. Here ‘characterized as’ is being used as a monster.
27. ‘In context C, I am speaking’ picks out me now matter who is speaking in context C. Why should it be any different for ‘looks green’?
28. The continuation is equally puzzling. At the end of Cohen’s story he says, ‘Mary and John agree that Smith doesn’t really know that the plane will stop in Chicago on the basis of the itinerary.’ The ‘know’ in that utterance should take on its semantic value in the Storytelling Context, and it’s peculiar how Mary and John end up agreeing with an utterance in the Storytelling Context.
29. To remind yourself of the justification for this modal claim see explanation under the Known Rupert Story earlier.
30. Here’s yet another attempt at making the same point. According to the contextualist about ‘tall’ it is context sensitive. But then how are we to make sense of the following seemingly coherent story: Everyone in my family is tall except my brother Bill, even though Bill is taller than my sister Mary. Still, my sister Mary is tall.

References

- Bezuidenhout, A., ‘Truth-Conditional Pragmatics’, *Philosophical Perspectives*, 16, 2002, pp. 105–134.
- Cappelen, H. and E. Lepore, ‘On an Alleged Connection between Indirect Quotation and Semantic Theory,’ *Mind and Language*, 12, 1997, pp. 278–296.
- , ‘Radical and Moderate Pragmatics’ (2003a).
- , ‘Perry on Unarticulated Constituents’, (2003b).
- , *Context and Sensitivity (book ms)*
- Castañeda, H-N, ‘The Theory of Questions, Epistemic Powers, and the Indexical Theory of Knowledge’, *Midwest Studies in Philosophy*, V (University of Minnesota Press), 1980, pp. 193–237.

- Cohen, Stuart, 'Contextualism, Skepticism, and the Structure of Reasons', *Philosophical Perspectives* 13, *Epistemology*, 1999, pp. 57–89.
- DeRose, Keith 'Contextualism and Knowledge Attributions,' *Philosophy and Phenomenological Research*, 52.4, 1992, pp. 913–929.
- , 'Solving the Skeptical Problem', in *The Philosophical Review* 104.1, 1995, pp. 1–52.
- , 'Single Scoreboard Semantics', unpublished ms (2002).
- Dreier, James, 'Internalism and Speaker Relativism', *Ethics*, 1990, pp. 6–26.
- Graff, Delia, 'Shifting Sands: An Interest Relative Theory of Vagueness', *Philosophical Topics*, 28.1, 2000, pp. 45–81.
- Kamp, Hans, 'Two Theories of Adjectives', in E. Keenan, ed. *Formal Semantics of Natural Language* (Cambridge: Cambridge University Press), 1975, pp. 23–55.
- Kaplan, David, 'Demonstratives.' in *Themes from Kaplan*. Oxford: Oxford University Press, J. Almog, J. Perry, and H. Wettstein 1989, pp. 481–563.
- Lepore, Ernie, 'Is there a problem with incomplete definite descriptions', *Descriptions*, eds. M. Reimer and A. Bezuidenhout, (Oxford University Press: Oxford), 2003, forthcoming.
- Lewis, David, 1979, 'Scorekeeping in a Language Game,' *Journal of Philosophical Logic* 8, 1979, pp. 339–359.
- , 'Index, Context, and Content', in *Philosophy and Grammar*, eds. S. Kanger and S. Ohman, (eds.), Dordrecht, Reidel, 1980; reprinted in *Papers in Philosophical Logic*, Cambridge University Press, 1998, pp. 21–44.
- , 'Elusive Knowledge', *Papers in Metaphysics and Epistemology* 74 *Australasian Journal of Philosophy*, 1996, pp. 549–67.
- Ludlow, Peter, 'Implicit Comparison Classes', *Linguistics and Philosophy*, 12, 1989, pp. 519–533.
- Moravcsik, J. M., *Meaning, Creativity, and the Partial Inscrutability of the Human Mind* (Vol. 79). Stanford: CSLI Publications, 1998.
- Parsons, T., 'Some Problems Concerning the Logic of Grammatical Modifiers', in Harman and Davidson (eds.), *Semantics of Natural Language* (Dordrecht, Reidel), 1972, pp. 127–41.
- Rysiew, Patrick, 'The Context-Sensitivity of Knowledge Attributions', *Noûs* 35:4, 2001, pp. 477–514.
- Searle, John, 'Literal Meaning', *Erkenntnis*, 13, 1978, pp. 207–224.
- , 'The Background of Meaning', in *Speech Act Theory and Pragmatics*, eds., Searle, J., F. Kiefer and M. Bierwisch (Reidel: Dordrecht), 1980, pp. 221–232.
- Soames, Scott, *Understanding Truth*, Oxford University Press, December 1998.
- Stanley, J. 'Context and Logical Form', *Linguistics and Philosophy*, 23, 2000, pp. 391–424.
- , 'Making it Articulated', *Mind and Language* 17, 2002, pp. 149–68.
- , 'On the Case for Contextualism', unpublished ms, 2003.
- Stanley, Jason and Szabo, Zoltan, 'On Quantifier Domain Restriction', *Mind and Language* 15.2, 2000, pp. 219–261.
- Travis, Charles, 'On What is Strictly Speaking True,' *Canadian Journal of Philosophy*, 15, 1985, pp. 187–229.
- , 'Meaning's Role in Truth', *Mind*, 100, 1996, pp. 451–466.
- , 'Reply to Simmons', *Mind*, 106, 1997, pp. 119–20.
- Unger, Peter, *Philosophical Relativity* (University of Minnesota Press), 1984.
- , 'Contextual Analysis in Ethics', *Philosophy and Phenomenological Research*, March 1995, pp. 1–26
- Westerståhl, Dag, 'Quantifiers in Formal and Natural Languages', in Gabbay, D. and Guenther, F. (ed.) *Handbook of Philosophical Logic*, Vol. IV, (Dordrecht: Kluwer), 1989, pp. 1–131.