

Cooperative Principle is being observed, and since there may be various possible specific explanations, a list of which may be open, the conversational implicatum in such cases will be disjunction of such specific explanations; and if the list of these is open, the implicatum will have just the kind of indeterminacy that many actual implicata do in fact seem to possess.

Further Notes on Logic and Conversation

I would like to begin by reformulating, in outline, the position which I took in Essay 2. I was operating, provisionally, with the idea that, for a large class of utterances, the total signification of an utterance may be regarded as divisible in two different ways. First, one may distinguish, within the total signification, between what is said (in a favored sense) and what is implicated; and second, one may distinguish between what is part of the conventional force (or meaning) of the utterance and what is not. This yields three possible elements—what is said, what is conventionally implicated, and what is nonconventionally implicated—though in a given case one or more of these elements may be lacking. For example, nothing may be said, though there is something which a speaker makes as if to say. Furthermore, what is nonconventionally implicated may be (or again may not be) conversationally implicated. I have suggested (1) that the Cooperative Principle and some subordinate maxims are standardly (though not invariably) observed by participants in a talk exchange and (2) that the assumptions required in order to maintain the supposition that they are being observed (or so far as is possible observed) either at the level of what is said—or failing that, at the level of what is implicated—are in systematic correspondence with non-conventional implicata of the conversational type.

Before proceeding further, I should like to make one supplementary remark. When I speak of the assumptions required in order to maintain the supposition that the Cooperative Principle and maxims are being observed on a given occasion, I am thinking of assumptions that are nontrivially required; I do not intend to include, for example,

an assumption to the effect that some particular maxim is being observed, or is thought of by the speaker as being observed. This seemingly natural restriction has an interesting consequence with regard to Moore's "paradox." On my account, it will not be true that when I say that *p*, I conversationally implicate that I believe that *p*; for to suppose that I believe that *p* (or rather think of myself as believing that *p*) is just to suppose that I am observing the first maxim of Quality on this occasion. I think that this consequence is intuitively acceptable; it is not a natural use of language to describe one who has said that *p* as having, for example, "implied," "indicated," or "suggested" that he believes that *p*; the natural thing to say is that he has expressed (or at least purported to express) the belief that *p*. He has of course committed himself, in a certain way, to its being the case that he believes that *p*, and while this commitment is not a case of saying that he believes that *p*, it is bound up, in a special way, with saying that *p*. The nature of the connection will, I hope, become apparent when I say something about the function of the indicative mood.

In response to Essay 2, I was given in informal discussion an example which seemed to me, as far as it went, to provide a welcome kind of support for the picture I have been presenting, in that it appeared to exhibit a kind of interaction between the members of my list of maxims which I had not foreseen. Suppose that it is generally known that New York and Boston were blacked out last night, and A asks B whether C saw a particular TV program last night. It will be conversationally unobjectionable for B, who knows that C was in New York, to reply, *No, he was in a blacked-out city*. B could have said that C was in New York, thereby providing a further piece of just possibly useful or interesting information, but in preferring the phrase *a blacked-out city* he was implicating (by the maxim prescribing relevance) a more appropriate piece of information, namely, why C was prevented from seeing the program. He could have provided both pieces of information by saying, e.g. *He was in New York, which was blacked out*, but the gain would have been insufficient to justify the additional conversational effort.

In suggesting, at the end of Essay 2, five features which conversational implicatures must possess, or might be expected to possess, I was not going so far as to suggest that it is possible, in terms of some or all of these features, to devise a decisive test to settle the question whether a conversational implicature is present or not—a test, that is

to say, to decide whether a given proposition *p*, which is normally part of the total signification of the utterance of a certain sentence, is on such occasions a conversational (or more generally a nonconventional) implicature of that utterance, or is, rather, an element in the conventional meaning of the sentence in question. (I express myself loosely, but, I hope, intelligibly.) Indeed I very much doubt whether the features mentioned can be made to provide any such knock-down test, though I am sure that at least some of them are useful as providing a more or less strong *prima facie* case in favor of the presence of a conversational implicature. But I would say that any such case would at least have to be supported by a demonstration of the way in which what is putatively implicated could have come to be implicated (by a derivation of it from conversational principles and other data); and even this may not be sufficient to provide a decisive distinction between conversational implicature and a case in which what was originally a conversational implicature has become conventionalized.

Let us look at two features in turn. First, nondetachability. It may be remembered that I said that a conversational implicature might be expected to exhibit a fairly high degree of nondetachability insofar as the implicature was carried because of what is said, and not by virtue of the manner of expression. The implicature is nondetachable insofar as it is not possible to find another way of saying the same thing (or approximately the same thing) which simply lacks the implicature. The implicature which attaches to the word *try* exhibits this feature. One would normally implicate that there was a failure, or some chance of failure, or that someone thinks or thought there to be some chance of failure, if one said *A tried to do x*; this implicature would also be carried if one said *A attempted to do x*, *A endeavored to do x*, or *A set himself to do x*.

This feature is not a necessary condition of the presence of a conversational implicature, partly because, as stated, it does not appear if the implicature depends on the manner in which what is said has been said, and it is also subject to the limitation that there may be no alternative way of saying what is said, or no way other than one which will introduce peculiarities of manner, such as by being artificial or long-winded.

Neither is it a sufficient condition, since the implicatures of utterances which carry presuppositions (if there are such things) (*He has left off beating his wife*) will not be detachable; and should a question arise whether a proposition implied by an utterance is entailed or

an assumption to the effect that some particular maxim is being observed, or is thought of by the speaker as being observed. This seemingly natural restriction has an interesting consequence with regard to Moore's "paradox." On my account, it will not be true that when I say that p , I conversationally implicate that I believe that p ; for to suppose that I believe that p (or rather think of myself as believing that p) is just to suppose that I am observing the first maxim of Quality on this occasion. I think that this consequence is intuitively acceptable; it is not a natural use of language to describe one who has said that p as having, for example, "implied," "indicated," or "suggested" that he believes that p ; the natural thing to say is that he has expressed (or at least purported to express) the belief that p . He has of course committed himself, in a certain way, to its being the case that he believes that p , and while this commitment is not a case of saying that he believes that p , it is bound up, in a special way, with saying that p . The nature of the connection will, I hope, become apparent when I say something about the function of the indicative mood.

In response to Essay 2, I was given in informal discussion an example which seemed to me, as far as it went, to provide a welcome kind of support for the picture I have been presenting, in that it appeared to exhibit a kind of interaction between the members of my list of maxims which I had not foreseen. Suppose that it is generally known that New York and Boston were blacked out last night, and A asks B whether C saw a particular TV program last night. It will be conversationally unobjectionable for B, who knows that C was in New York, to reply, *No, he was in a blacked-out city*. B could have said that C was in New York, thereby providing a further piece of just possibly useful or interesting information, but in preferring the phrase *a blacked-out city* he was implicating (by the maxim prescribing relevance) a more appropriate piece of information, namely, why C was prevented from seeing the program. He could have provided both pieces of information by saying, e.g. *He was in New York, which was blacked out*, but the gain would have been insufficient to justify the additional conversational effort.

In suggesting, at the end of Essay 2, five features which conversational implicatures must possess, or might be expected to possess, I was not going so far as to suggest that it is possible, in terms of some or all of these features, to devise a decisive test to settle the question whether a conversational implicature is present or not—a test, that is

to say, to decide whether a given proposition p , which is normally part of the total signification of the utterance of a certain sentence, is on such occasions a conversational (or more generally a nonconventional) implicature of that utterance, or is, rather, an element in the conventional meaning of the sentence in question. (I express myself loosely, but, I hope, intelligibly.) Indeed I very much doubt whether the features mentioned can be made to provide any such knock-down test, though I am sure that at least some of them are useful as providing a more or less strong *prima facie* case in favor of the presence of a conversational implicature. But I would say that any such case would at least have to be supported by a demonstration of the way in which what is putatively implicated could have come to be implicated (by a derivation of it from conversational principles and other data); and even this may not be sufficient to provide a decisive distinction between conversational implicature and a case in which what was originally a conversational implicature has become conventionalized.

Let us look at two features in turn. First, nondetachability. It may be remembered that I said that a conversational implicature might be expected to exhibit a fairly high degree of nondetachability insofar as the implicature was carried because of what is said, and not by virtue of the manner of expression. The implicature is nondetachable insofar as it is not possible to find another way of saying the same thing (or approximately the same thing) which simply lacks the implicature. The implicature which attaches to the word *try* exhibits this feature. One would normally implicate that there was a failure, or some chance of failure, or that someone thinks or thought there to be some chance of failure, if one said *A tried to do x*; this implicature would also be carried if one said *A attempted to do x*, *A endeavored to do x*, or *A set himself to do x*.

This feature is not a necessary condition of the presence of a conversational implicature, partly because, as stated, it does not appear if the implicature depends on the manner in which what is said has been said, and it is also subject to the limitation that there may be no alternative way of saying what is said, or no way other than one which will introduce peculiarities of manner, such as by being artificial or long-winded.

Neither is it a sufficient condition, since the implicatures of utterances which carry presuppositions (if there are such things) (*He has left off beating his wife*) will not be detachable; and should a question arise whether a proposition implied by an utterance is entailed or

conversationally implicated, in either case the implication will be non-detachable. Reliance on this feature is effective primarily for distinguishing between certain conventional implicatures and nonconventional implicatures.

Second, cancelability. You will remember that a putative conversational implicature that *p* is explicitly cancelable if, to the form of words the utterance of which putatively implicates that *p*, it is admissible to add *but not p*, or *I do not mean to imply that p*, and it is contextually cancelable if one can find situations in which the utterance of the form of words would simply not carry the implicature. Now I think that all conversational implicatures are cancelable, but unfortunately one cannot regard the fulfillment of a cancelability test as decisively establishing the presence of a conversational implicature. One way in which the test may fail is connected with the possibility of using a word or form of words in a loose or relaxed way. Suppose that two people are considering the purchase of a tie which both of them know to be medium green; they look at it in different lights, and say such things as *It is a light green now*, or *It has a touch of blue in it in this light*. Strictly (perhaps) it would be correct for them to say *It looks light green now* or *It seems to have a touch of blue in it in this light*, but it would be unnecessary to put in such qualificatory words, since both know (and know that the other knows) that there is no question of a real change of color. A similar linguistic phenomenon attends such words as *see*: If we all know that Macbeth hallucinated, we can quite safely say that Macbeth saw Banquo, even though Banquo was not there to be seen, and we should not conclude from this that an implication of the existence of the object said to be seen is not part of the conventional meaning of the word *see*, nor even (as some have done) that there is one sense of the word *see* which lacks this implication.

Let us consider this point in relation to the word *or*. Suppose that someone were to suggest that the word *or* has a single "strong" sense, which is such that it is part of the meaning of *A or B* to say (or imply) not only (1) that $A \vee B$, but also (2) that there is some non-truth-functional reason for accepting that $A \vee B$, i.e. that there is some reasonable (though not necessarily conclusive) argument with $A \vee B$ as conclusion which does not contain one of the disjuncts as a step (does not proceed via *A* or via *B*). Now it would be easy to show that the second of the two suggested conditions is cancelable: I can say to my children at some stage in a treasure hunt, *The prize is either in the*

garden or in the attic. I know that because I know where I put it, but I'm not going to tell you. Or I could just say (in the same situation) *The prize is either in the garden or in the attic*, and the situation would be sufficient to apprise the children of the fact that my reason for accepting the disjunction is that I know a particular disjunct to be true. And in neither case would I be implying that there is a non-truth-functional ground, though I am not relying on it; very likely there would not be such a ground. To this objection, the "strong" theorist (about *or*) might try the move "Ah, but when you say $A \vee B$, without meaning to imply the existence of a non-truth-functional ground, you are using $A \vee B$ loosely, in a relaxed way which the nature of the context of utterance makes permissible." At this point, we might either (1) produce further cancellation cases, which were less amenable to representation as "loose" uses, for example, to the appearance of disjunctions as the antecedents of conditionals (*If the prize is either in the garden or in the attic, Johnny will find it first*), or (2) point out that to characterize a use as "loose" carries certain consequences which are unwelcome in this case—if to say *Macbeth saw Banquo* is to speak loosely, then I speak "under license" from other participants; if someone objects, there is at least some onus on me to speak more strictly. But not even a stickler for correct speech could complain about the utterance (in the described circumstances) of *The prize is either in the garden or in the attic*.

But the strong theorist has another obvious resource. He may say that there are two senses of the word *or*, a strong one and a weak (truth-functional) one, and that all that is shown by the success of the cancelability test is that here the sense employed was the weak one. To counter this suggestion, we might proceed in one or more of the following ways.

1. We might argue that if *or* is to be supposed to possess a strong sense, then it should be possible to suppose it (*or*) to bear this sense in a reasonably wide range of linguistic settings; it ought to be possible, for example, to say *It is not the case that A or B* or *Suppose that A or B*, where what we are denying, or inviting someone to suppose, is that $A \vee B$ (in the strong sense of *or*). But this, in the examples mentioned, does not seem to be possible; in anything but perhaps a very special case to say *It is not the case that A or B* seems to amount to saying that neither *A* nor *B* (that is, cannot be interpreted as based on a denial of the second condition), and to say *Suppose that A or B* seems to amount to inviting someone to suppose merely that

one of the two disjuncts is true. A putative second sense of *or* should not be so restricted in regard to linguistic setting as that, and in particular should not be restricted to “unenclosed” occurrences of *A or B*—for these an alternative account (in terms of implicature) is readily available. The strong theorist might meet a part of this attack by holding that the second condition is not to be thought of as part of what is said (or entailed) by saying *A or B*, and so not as something the denial of which would justify the denial of *A or B*; it should rather be thought of as something which is conventionally implicated. And to deny *A or B* might be to implicate that there was some ground for accepting *A or B*. But he is then open to the reply that, if a model case for a word which carries a conventional implicature is *but*, then the negative form *It is not the case that A or B*, if to be thought of as involving *or* in the strong sense, should be an uncomfortable thing to say, since *It is not the case that A but B* is uncomfortable. In any case the nature of conventional implicature needs to be examined before any free use of it, for explanatory purposes, can be indulged in.

2. We might try to convince the strong theorist that if *or* is to be regarded as possessing a strong sense as well as a weak one, the strong sense should be regarded as derivative from the weak one. The support for this contention would have to be a combination of two points: (a) that the most natural expression of the second condition involves a use of *or* in the “weak” sense; and even if the weak use of *or* is avoided the idea seems to be explicitly involved; it is difficult to suppose that people could use a word so as to include in its meaning that there is evidence of a certain sort for a proposition without having a distinct notion of that for which the evidence is evidence. (b) One who says that *A or B*, using *or* truth-functionally, could be shown in normal circumstances to implicate (conversationally) that there are non-truth-functional grounds for supposing that $A \vee B$. For to say that *A or B* (interpreted weakly) would be to make a weaker and so less informative statement than to say that *A* or to say that *B*, and (on the assumption, which I shall not here try to justify, that it would be of interest to an audience to know that one of the disjuncts is true) would therefore be to make a less informative statement than would be appropriate in the circumstances. So there is an implicature (provided the speaker is not opting out) that he is not in a position to make a stronger statement, and if, in conformity with the second maxim of Quality, the speaker is to be presumed to have evidence for what he says, then the speaker thinks that there are non-truth-

functional grounds for accepting *A or B*. We might next argue that if the strong sense of *or* is derivative from the weaker sense, then it ought to conform to whatever general principles there may be which govern the generation of derivative senses. This point is particularly strong in connection with a suggestion that *or* possesses a derivative sense; for we are not particularly at home with the application of notions such as “meaning” and “sense” to words so nondescriptive as *or*; the difficulties we encounter here are perhaps similar to, though not so severe as, the difficulties we should encounter if asked to specify the meaning or meanings of a preposition like *to* or *in*. So I suspect that we should need to rely fairly heavily on an application to the case of *or* of whatever general principles there may be which apply to more straightforward cases and which help to determine when a derivative sense should be supposed to exist, and when it should not.

It might be objected that whether one sense of a word is to be regarded as derivative from another sense of that word should be treated as a question about the history of the language to which the word belongs. This may be so in general (though in many cases it is obvious, without historical research, that one sense must be secondary to another), but if I am right in thinking that conversational principles would not allow the word *or* to be used in normal circumstances without at least an implicature of the existence of non-truth-functional grounds, then it is difficult to see that research could contribute any information about temporal priority in this case.

I offer three further reflections about the proliferation of senses.

1. I would like to propose for acceptance a principle which I might call Modified Occam’s Razor, *Senses are not to be multiplied beyond necessity*. Like many regulative principles, it would be a near platitude, and all would depend on what was counted as “necessity.” Still, like other regulative principles, it may guide. I can think of other possible precepts which would amount to much the same. One might think, for example, of not allowing the supposition that a word has a further (and derivative) sense unless the supposition that there is such a sense does some work, explains why our understanding of a particular range of applications of the word is so easy or so sure, or accounts for the fact that some application of the word outside that range, which would have some prima facie claim to legitimacy, is in fact uncomfortable. Again one might formulate essentially the same idea by recommending that one should not suppose what a speaker would mean when he used a word in a certain range of cases to count

as a special sense of the word, if it should be predictable, independently of any supposition that there is such a sense, that he would use the word (or the sentence containing it) with just that meaning. If one makes the further assumption that it is more generally feasible to strengthen one's meaning by achieving a superimposed implicature, than to make a relaxed use of an expression (and I don't know how this assumption would be justified), then Modified Occam's Razor would bring in its train the principle that one should suppose a word to have a less restrictive rather than a more restrictive meaning, where choice is possible.

What support would there be for Modified Occam's Razor? Perhaps we might look at two types of example of real or putative derivative senses. One type (unlike the case of *or*) would involve "transferred" senses; the other would involve derivative senses which are specificatory of the original senses (the proposed derivative sense of "or" would be a special case of this kind).

a. Consider such adjectives as *loose*, *unfettered*, and *unbridled* in relation to a possible application to the noun *life*. (I assume that such an application of each word would not be nonderivative or literal; that the ambiguous expression *a loose liver* would involve a nonderivative sense of *loose* if uttered for example by a nurse in a hospital who complained about the number of patients with loose livers, but not if uttered censoriously to describe a particular man.) It seems to me that (in the absence of any further sense for either word) one might expect to be able to mean more or less the same by *a loose life*, and *an unfettered life*; the fact that, as things are, *loose life* is tied to dissipation, whereas *unfettered life* seems quite general in meaning, suggests that perhaps *loose* does, and *unfettered* does not, have a derivative sense in this area. As for *unbridled life* (which one might perhaps have expected, *prima facie*, to mean much the same as *unfettered life*), the phrase is slightly uncomfortable (because *unbridled* seems to be tied to such words as *passion*, *temper*, *lust*, and so on).

b. As for words with specificatory derivative senses, there seems to be some tendency for one of two things to happen: Either the original general sense becomes obsolete (like *car*, meaning "wheeled vehicle"), or the specificatory condition takes over; we should perhaps continue to call gramophone records *discs* even if (say) they came to be made square (provided they remain not too unlike discs, in the original sense of the word), and perhaps the word *cylinder* exemplifies the same feature. But there are words of which neither is true: an obvious example is the word *animal* (meaning (i) "member of animal king-

dom," (ii) "beast"). There is here some sort of a parallel, in relation to Modified Occam's Razor and its variants, between *animal* and the candidate word *or*. *Animal* perhaps infringes a weak principle to the effect that a further sense should not be recognized if, *on the assumption* that the word were to have a specificatory further sense, the identity of that sense would be predictable; for it could no doubt be predicted that *if* the word *animal* were to have such a sense, it would be one in which the word did not apply to human beings. But it would seem not to be predictable (history of language apart) that anyone would *in fact* use the word *animal* to mean "beast," whereas given a truth-functional *or* it is predictable (assuming conversational principles) that people would use *A or B* to imply the existence of non-truth-functional grounds. So, at least, so far as I can see (not far, I think), there is as yet no reason not to accept Modified Occam's Razor.

2. We must of course give due (but not undue) weight to intuitions about the existence or nonexistence of putative senses of a word (how could we do without them?). Indeed if the scheme which I have been proposing is even proceeding in the right direction, at least some reliance must be placed on such intuitions. For in order that a nonconventional implicature should be present in a given case, my account requires that a speaker shall be able to utilize the conventional meaning of a sentence. If nonconventional implicature is built on what is said, if what is said is closely related to the conventional force of the words used, and if the presence of the implicature depends on the intentions of the speaker, or at least on his assumptions, with regard to the possibility of the nature of the implicature being worked out, then it would appear that the speaker must (in some sense or other of the word *know*) know what is the conventional force of the words which he is using. This indeed seems to lead to a sort of paradox: If we, as speakers, have the requisite knowledge of the conventional meaning of sentences we employ to implicate, when uttering them, something the implication of which depends on the conventional meaning in question, how can we, as theorists, have difficulty with respect to just those cases in deciding where conventional meaning ends and implicature begins? If it is true, for example, that one who says *that A or B* implicates the existence of non-truth-functional grounds for *A or B*, how can there be any doubt whether the word "or" has a strong or weak sense? I hope that I can provide the answer to this question, but I am not certain that I can.

3. I have briefly mentioned a further consideration bearing on the

question of admissibility of a putative sense of a word, namely, whether on the supposition that the word has that sense, there would be an adequate range of linguistic environments in which the word could be supposed to bear that sense. Failure in this respect would indicate an implicature or an idiom.

There are, I am certain, other possible principles which ought to be considered with regard to the proliferation of senses. In particular I have said nothing explicitly about the adequacy of substitutability tests. But I propose to leave this particular topic at this point.

I have so far been considering two questions. (1) On the assumption that a word has only one conventional meaning (or only one relevant conventional meaning), how much are we to suppose to be included in that meaning? (2) On the assumption that a word has at least one conventional meaning (or relevant conventional meaning), are we to say that it has one, or more than one, such meaning? In particular, are we to ascribe to it a second sense or meaning, derivative from or dependent on a given first meaning or sense? We should consider also examples of elements in or aspects of utterances which, not being words, are candidates for conventional meaning (or significance). These include stress, irony, and truth.

Stress

Some cases of stress are clearly relevant to the possession of conventional meaning, such as (fixed) stress on particular syllables or a word, as in the contrast between "c^ontent" with the stress on the "o" and c^ont^ent with the stress on "tent"; we should not, of course, here assign meaning to the stress itself. I am concerned not with cases of that sort, but with cases in which we think of a word as being stressed, and variably so, stressed on some occasions but not on others.

We might start by trying to think of stress as a purely natural way of highlighting, or making prominent, a particular word; one might compare putting some object (such as a new hat) in an obvious place in a room so that someone coming into the room will notice or pay attention to it. But there are various suggestible ways of doing this with a word, such as intoning it or saying it in a squeaky voice. Such methods would not just be thought unusual, they would be frowned on. They would also very likely fail to achieve the effect of highlighting just because there is an approved way of doing this. So there is a

good case for regarding stress as a conventional device for highlighting. But to say this much is not to assign to stress a conventional significance or meaning: it is only to treat it as a conventional way of fulfilling a certain purpose, which is not yet established as a purpose connected with communication. But stress clearly does in fact on many occasions make a difference to the speaker's meaning; indeed it is one of the elements which help to generate implicatures. Does this fact require us to attribute any conventional meaning to stress?

In accordance with the spirit of Modified Occam's Razor, we might attribute conventional meaning to stress only if it is unavoidable. Thus we might first introduce a slight extension to the maxim enjoining relevance, making it apply not only to what is said but also to features of the means used for saying what is said. This extension will perhaps entitle us to expect that an aspect of an utterance which it is within the power of the speaker to eliminate or vary, even if it is introduced unreflectively, will have a purpose connected with what is currently being communicated; unless, of course, its presence can be explained in some other way.

At least three types of context in which stress occurs seem to invite ordering:

(1) One such context includes replies to W-questions ("who," "what," "why," "when," and "where"):

A: *Who paid the bill?* B: *Jónes did.*

A: *What did Jones do to the cat?* B: *He kⁱcked it.*

It also includes exchanges of such forms as:

A: $S(\alpha)$

B: $\bar{S}(\acute{\alpha}): S(\acute{\beta})$

For example:

A: *Jones paid the bill.*

B: *Jónes didn't pay the bill; Smíth paid it.*

In such examples, stress is automatic or a matter of habit (maybe difficult to avoid), and we are not inclined to say that anything is meant or implicated. However, the effect is to make perspicuous elements which complete open sentences for which questions (in effect) demand completion, or elements in respect of which what B is prepared to assert (or otherwise say) and what B has asserted differ.

(2) Another context in which stress occurs includes such cases as

incomplete versions of the conversational schema exemplified in the preceding example. Without a preceding statement to the effect that Jones paid the bill, B says *Jónes didn't pay the bill; Smith did*. Here, given that this sentence is to be uttered, the stress may be automatic, but the remark is not prompted by a previous remark (it is volunteered), and we are inclined to say that the implicature is that someone thinks or might think that Jones did pay the bill. The maxim of relation requires that B's remark should be relevant to something or other, and B, by speaking as he would speak in reply to a statement that Jones paid the bill, shows that he has such a statement in mind.

(3) In a third context B just says *Jónes didn't pay the bill*. B speaks as if he were about to continue as in the previous context. B therefore implicates that someone (other than Jones) paid the bill.

In general, $S(\alpha)$ is contrasted with the result of substituting some expression β for α , and commonly the speaker suggests that he would deny the substitute version, but there are other possibilities. For example, *I knéw that* may be contrasted with *I believed that*, and the speaker may implicate not that he would deny *I believed that p*, but that he would not confine himself to such a weaker statement (with the implicit completion *I didn't merely believe it*).

This last point has relevance to the theory of "knowledge." According to a certain "strong" account of knowledge:

- A knows that p just in case*
- (1) *p*
 - (2) *A thinks that p*
 - (3) *A has conclusive evidence that p*

This presents possible difficulties of a regressive nature:

- (1) Does A have to know that the evidence for *p* is true?
- (2) Does A have to know that the evidence is conclusive?

But in general the theory seems *too* strong. An examination candidate at an oral knows the date of the battle of Waterloo. He may know this without conclusive evidence; he may even answer after hesitation (showing in the end that he knows the answer). I suggest something more like the following:

- A knows that p just in case*
- (1) *p*
 - (2) *A thinks that p*

- (3) *Some conditions placing restriction on how he came to think p* (cf. causal theory)

If I say *I know that p*, then perhaps sometimes there is a nonconventional implicature of strong or conclusive evidence (not mere thinking that *p*, with *p* true)—cf. *He loves her*. And this is not the only interpretation of stress; it can also mean, "You don't need to tell me."

Irony

The second example of an element in, or aspect of, some utterances, with regard to which there might be some doubt whether or not it has a conventional meaning, emerges from my (too) brief characterization of irony in Essay 2. Discussion with Rogers Albritton showed me that something is missing in this account. It seems very dubious whether A's knowledge that B has been cheated by C, that B knows that A knows that this is so, that B's remark *He is a fine friend* is to be presumed to relate to this episode, and that the remark is seemingly false (even obviously false), is enough to ensure, with reasonable certainty, that A will suppose B to mean the negation of what he has made as if to say. A might just be baffled, or might suppose that, despite the apparent falsity of the remark, B was meaning something like *He is, usually, a fine friend: how could he have treated me like that?* It was suggested to me that what should have been mentioned in my account was, first, a familiarity with the practice of using a sentence, which would standardly mean that *p*, in order to convey that not-*p* (a familiarity which might be connected with a natural tendency in us to use sentences in this way), and, second, an ironical tone in which such utterances are made, and which (perhaps) conventionally signifies that they are to be taken in reverse.

This suggestion does not seem to me to remedy the difficulty. Consider the following example. A and B are walking down the street, and they both see a car with a shattered window. B says, *Look, that car has all its windows intact*. A is baffled. B says, *You didn't catch on; I was in an ironical way drawing your attention to the broken window*. The absurdity of this exchange is I think to be explained by the fact that irony is intimately connected with the expression of a feeling, attitude, or evaluation. I cannot say something ironically un-

less what I say is intended to reflect a hostile or derogatory judgment or a feeling such as indignation or contempt. I can for example say *What a scoundrel you are!* when I am well disposed toward you, but to say that will be playful, not ironical, and will be inappropriate unless there is some shadow of justification for a straightforward application—for example you have done something which some people (though not I) might frown upon. If when you have just performed some conspicuously disinterested action I say, *What an egotist you are, always giving yourself the satisfaction of doing things for other people!* I am expressing something like what might be the reaction of an extreme cynic. Whereas to say *He's a fine friend* is unlikely to involve any hint of anyone's approval.

I am also doubtful whether the suggested vehicle of signification, the ironical tone, exists as a specific tone; I suspect that an ironical tone is always a contemptuous tone, or an amused tone, or some other tone connected with one or more particular feelings or attitudes; what qualifies such a tone as ironical is that it appears, on this and other occasions, when an ironical remark is made. This question could no doubt be settled by experiment. Even if, however, there is no specifically ironical tone, it still might be suggested that a contemptuous or amused tone, when conjoined with a remark which is blatantly false, conventionally indicates that the remark is to be taken in reverse. But the suggestion does not seem to me to have much plausibility. While I may without any inappropriateness prefix the employment of a metaphor with *to speak metaphorically*, there would be something very strange about saying, *to speak ironically, he is a splendid fellow*. To be ironical is, among other things, to pretend (as the etymology suggests), and while one wants the pretense to be recognized as such, to announce it as a pretense would spoil the effect. What is possibly more important, it might well be essential to an element's having conventional significance that it could have been the case that some quite different element should have fulfilled the same semantic purpose; that if a contemptuous tone does conventionally signify in context that a remark is to be taken in reverse, then it might have been, for example, that a querulous tone should have been used (instead) for the same purpose. But the connection of irony with the expression of feeling seems to preclude this; if speaking ironically has to be, or at least appear to be, the expression of a certain feeling or attitude, then a tone suitable to such a feeling or attitude seems to be mandatory, at least for the unsophisticated examples.

Truth

Among the A-philosophical theses which I considered in Essay 1 was the original version of a "speech-act" account of truth, proposed by Strawson many years ago¹ though extensively modified by him since then. He was influenced, I think, by four main considerations: (1) that the word *true* is properly, or at least primarily, to be applied to statements (what is stated), in view of the difficulties which he thought he saw in the thesis that it should be understood as applying to utterances; (2) that given the correctness of the previous supposition, no theory which treats truth as consisting in a relation (or correlation) between statements and facts is satisfactory, since statements and facts cannot be allowed to be distinct items in the real world; (3) that Ramsey's account of truth²—namely, that to assert that a proposition is true is to assert that proposition—is correct so far as it goes; and (4) that it does not go far enough, since it omits to take seriously the fact that we should not always be willing to tolerate the substitution of, for example, *It is true that it is raining* for *It is raining*. So he propounded the thesis that to say of a statement that it is true is (1) insofar as it is to assert anything, to assert that statement and (2) not merely to assert it but to endorse, confirm, concede, or reassert it (the list is not, of course, intended to be complete).

Such a theory seems to me to have at least two unattractive features, on the assumption that it was intended to give an account of the meaning (conventional significance) of the word *true*. (1) (A familiar type of objection) it gives no account, or no satisfactory account, of the meaning of the word *true* when it occurs in unasserted subsentences (e.g. *He thinks that it is true that . . .* or *If it is true that . . .*). (2) It is open to an objection which I am inclined to think holds against Ramsey's view (of which the speech-act theory is an offshoot). A theory of truth has (as Tarski noted) to provide not only for occurrences of *true* in sentences in which what is being spoken of as true is specified, but also for occurrences in sentences in which no specification is given (e.g. *The policeman's statement was true*). According both to the speech-act theory, I presume, and to Ramsey's theory, at least part of what the utterer of such a sentence is doing is to assert whatever it was that the policeman stated. But the utterer may not

1. P. F. Strawson, "Truth," *Analysis* 9, no. 6 (1949).

2. *Foundations of Mathematics*, pp. 142–143.

know what that statement was; he may think that the policeman's statement was true because policemen always speak the truth, or that that policeman always speaks the truth, or that policeman in those circumstances could not but have spoken the truth. Now assertion presumably involves committing oneself, and while it is possible to commit oneself to a statement which one has not identified (I could commit myself to the contents of the Thirty-Nine Articles of the Church of England, without knowing what they say), I do not think I should be properly regarded as having committed myself to the content of the policeman's statement, merely in virtue of having said that it was true. When to my surprise I learn that the policeman actually said, *Monkeys can talk*, I say (perhaps), *Well, I was wrong*, not *I withdraw that*, or *I withdraw my commitment to that*. I never was committed to it.

My sympathies lie with theories of the correspondence family, which Strawson did (and I think still does) reject, but it is not to my present purpose (nor within my capacities) to develop adequately any such theory. What I wish to do is to show that, on the assumption that a certain sort of theory of this kind is correct, then, with the aid of the apparatus discussed in Essay 2 it is possible to accommodate the linguistic phenomena which led Strawson to formulate the original version of the speech-act theory. Let me assume (and hope) that it is possible to construct a theory which treats truth as (primarily) a property of utterances; to avoid confusion I shall use, to name such a property, not "true" but "factually satisfactory." Let me also assume that it will be a consequence of such a theory that there will be a class *K* of utterances (utterances of affirmative subject-predicate sentences) such that every member of *K* (1) *designates*³ some item and *indicates*³ some class (these verbs to be explained within the theory), and (2) is factually satisfactory if the item belongs to the class. Let me finally assume that there can be a method of introducing a form of expression *It is true that . . .* and linking it with the notion "factually satisfactory," a consequence of which will be that to say *It is true that Smith is happy* will be equivalent to saying that any utterance of class *K* which designates Smith and indicates the class of happy people is factually satisfactory (that is, any utterance which assigns Smith to the class of happy people is factually satisfactory).

If some such account of *It is true that . . .* is correct (or indeed any

account which represents saying *It is true that p* as equivalent to saying something about utterances) then it is possible to deal with the linguistic facts noted by Strawson. To say *Smith is happy* is not to make a (concealed) reference to utterances of a certain sort, whereas to say *It is true that Smith is happy* is to do just that, though of course if Smith is happy, it is true that Smith is happy. If I choose the form which does make a concealed reference to utterances, and which is also the more complex form, in preference to the simpler form, it is natural to suppose that I do so because an utterance to the effect that Smith is happy has been made by myself or someone else, or might be so made. Such speech acts as endorsing, agreeing, confirming, and conceding, which Strawson (presumably) supposed to be conventionally signaled by the use of the word *true* are just those which, in saying in response to some remark "That's true," one would be performing (without any special signal). And supposing no one actually to have said that Smith is happy, if I say "It is true that Smith is happy" (e.g. concessively) I shall implicate that someone might say so; and I do not select this form of words as, for example, a response to an inquiry whether Smith is happy when I do not wish this implicature to be present.

3. These verbs to be explained within the theory.